

q3

February 20, 2024

```
[47]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
import scipy.stats as stats

def season_category(x):
    if x == 1:
        return 'season_1'
    elif x == 2:
        return 'season_2'
    elif x == 3:
        return 'season_3'
    else:
        return 'season_4'
```

Reading the dataset

```
[48]: df = pd.read_csv('BRSM_Assignment_Datasets.csv')
data = df
print(df.head())
print()
print("Columns are given by:")
print(df.columns)
alpha = 0.05
```

	datetime	season	holiday	workingday	weather	temp	atemp	\
0	2011-01-01 00:00:00	1	0	0	1	9.84	14.395	
1	2011-01-01 01:00:00	1	0	0	1	9.02	13.635	
2	2011-01-01 02:00:00	1	0	0	1	9.02	13.635	
3	2011-01-01 03:00:00	1	0	0	1	9.84	14.395	
4	2011-01-01 04:00:00	1	0	0	1	9.84	14.395	

	humidity	windspeed	casual	registered	count
0	81	0.0	3	13	16
1	80	0.0	8	32	40
2	80	0.0	5	27	32
3	75	0.0	3	10	13

```
4          75          0.0          0          1          1
```

Columns are given by:

```
Index(['datetime', 'season', 'holiday', 'workingday', 'weather', 'temp',  
      'atemp', 'humidity', 'windspeed', 'casual', 'registered', 'count'],  
      dtype='object')
```

Shape of the dataset

```
[49]: df.shape
```

```
[49]: (10886, 12)
```

Converting the datatype of datetime column from object to datetime

```
[50]: df['datetime'] = pd.to_datetime(df['datetime'])
```

```
[51]: df['season'] = df['season'].apply(season_category)
```

```
[52]: df['season'] = df['season'].astype('category')  
df['holiday'] = df['holiday'].astype('category')  
df['workingday'] = df['workingday'].astype('category')  
df['weather'] = df['weather'].astype('category')  
df['temp'] = df['temp'].astype('float32')  
df['atemp'] = df['atemp'].astype('float32')  
df['humidity'] = df['humidity'].astype('float32')  
df['windspeed'] = df['windspeed'].astype('float32')  
df['casual'] = df['casual'].astype('int32')  
df['registered'] = df['registered'].astype('int32')  
df['count'] = df['count'].astype('int32')
```

```
[53]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 10886 entries, 0 to 10885  
Data columns (total 12 columns):  
#   Column          Non-Null Count  Dtype  
---  -  
0   datetime        10886 non-null  datetime64[ns]  
1   season          10886 non-null  category  
2   holiday         10886 non-null  category  
3   workingday      10886 non-null  category  
4   weather         10886 non-null  category  
5   temp            10886 non-null  float32  
6   atemp           10886 non-null  float32  
7   humidity        10886 non-null  float32  
8   windspeed       10886 non-null  float32  
9   casual          10886 non-null  int32
```

```

10 registered 10886 non-null int32
11 count      10886 non-null int32
dtypes: category(4), datetime64[ns](1), float32(4), int32(3)
memory usage: 426.0 KB

```

```
[54]: df.describe()
```

```
[54]:
```

	datetime	temp	atemp	\
count	10886	10886.000000	10886.000000	
mean	2011-12-27 05:56:22.399411968	20.230862	23.655085	
min	2011-01-01 00:00:00	0.820000	0.760000	
25%	2011-07-02 07:15:00	13.940000	16.665001	
50%	2012-01-01 20:30:00	20.500000	24.240000	
75%	2012-07-01 12:45:00	26.240000	31.059999	
max	2012-12-19 23:00:00	41.000000	45.455002	
std	NaN	7.791590	8.474601	

	humidity	windspeed	casual	registered	count
count	10886.000000	10886.000000	10886.000000	10886.000000	10886.000000
mean	61.886459	12.799396	36.021955	155.552177	191.574132
min	0.000000	0.000000	0.000000	0.000000	1.000000
25%	47.000000	7.001500	4.000000	36.000000	42.000000
50%	62.000000	12.998000	17.000000	118.000000	145.000000
75%	77.000000	16.997900	49.000000	222.000000	284.000000
max	100.000000	56.996899	367.000000	886.000000	977.000000
std	19.245033	8.164537	49.960477	151.039033	181.144454

```
[55]: def plot_categorical_distribution(df, column, subplot_index):
    column_distribution = df[column].value_counts().reset_index()
    column_distribution.columns = [column, 'count']
    plt.subplot(subplot_index)
    plt.pie(column_distribution['count'], labels=column_distribution[column],
    autopct='%1.1f%%', startangle=140)
    plt.title(f'Distribution of {column}')
    plt.axis('equal')

plt.figure(figsize=(10, 10))

plt.subplot(2, 2, 1)
plot_categorical_distribution(df, 'season', 221)


plt.subplot(2, 2, 2)
plot_categorical_distribution(df, 'weather', 222)

plt.subplot(2, 2, 3)
plot_categorical_distribution(df, 'holiday', 223)

```

```
plt.subplot(2, 2, 4)
plot_categorical_distribution(df, 'workingday', 224)

plt.tight_layout()
plt.show()
```



```
[56]: def plot_countplot(df, column, subplot_index):
    plt.subplot(subplot_index)
    sns.countplot(data=df, x=column)
    plt.title(f'Countplot of {column}')

plt.figure(figsize=(12, 10))
```

```

plt.subplot(2, 2, 1)
plot_countplot(df, 'season', 221)

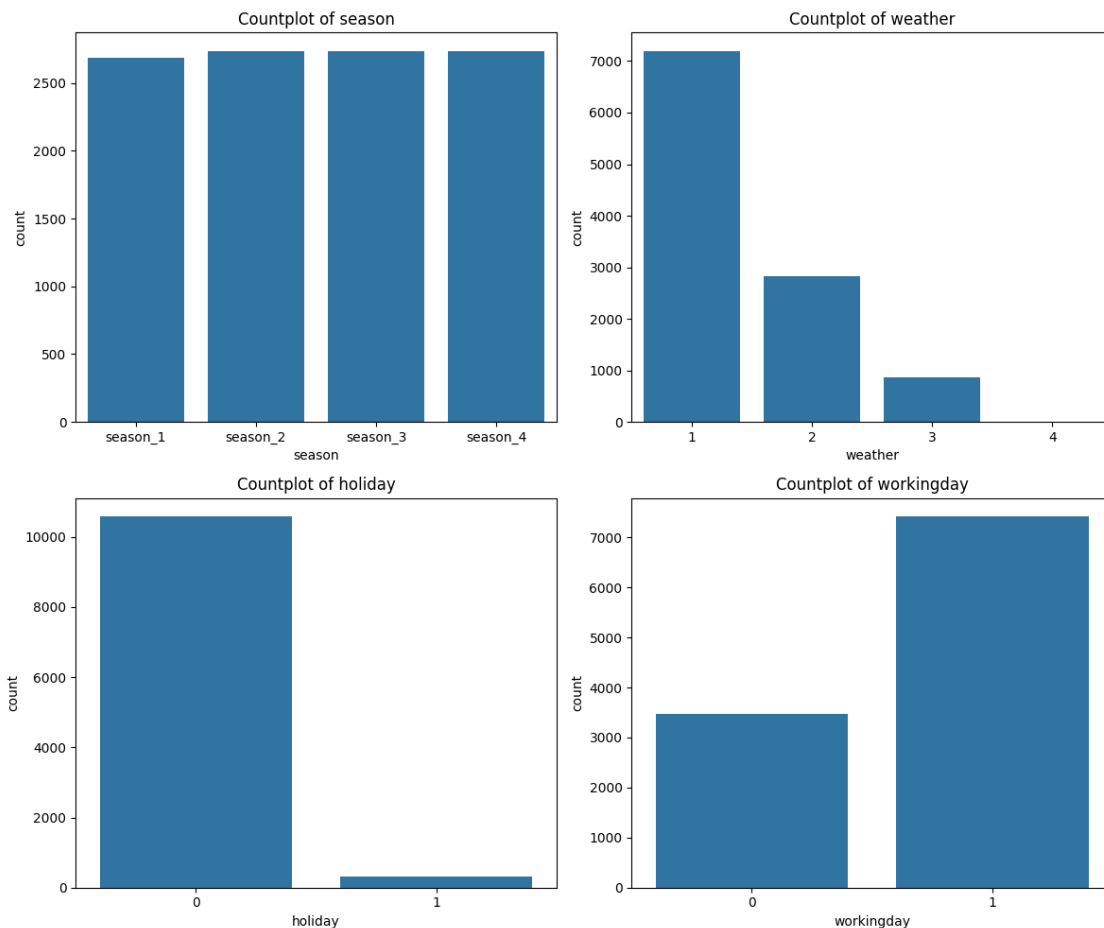
plt.subplot(2, 2, 2)
plot_countplot(df, 'weather', 222)

plt.subplot(2, 2, 3)
plot_countplot(df, 'holiday', 223)

plt.subplot(2, 2, 4)
plot_countplot(df, 'workingday', 224)

plt.tight_layout()
plt.show()

```



```

[57]: def plot_histplot(df, column, subplot_index):
      plt.subplot(subplot_index)
      sns.histplot(data=df, x=column, kde=True, bins=50)

```

```

plt.title(f'Histogram of {column}')

plt.figure(figsize=(18, 12))

plt.subplot(2, 3, 1)
plot_histplot(df, 'temp', 231)

plt.subplot(2, 3, 2)
plot_histplot(df, 'atemp', 232)

plt.subplot(2, 3, 3)
plot_histplot(df, 'humidity', 233)

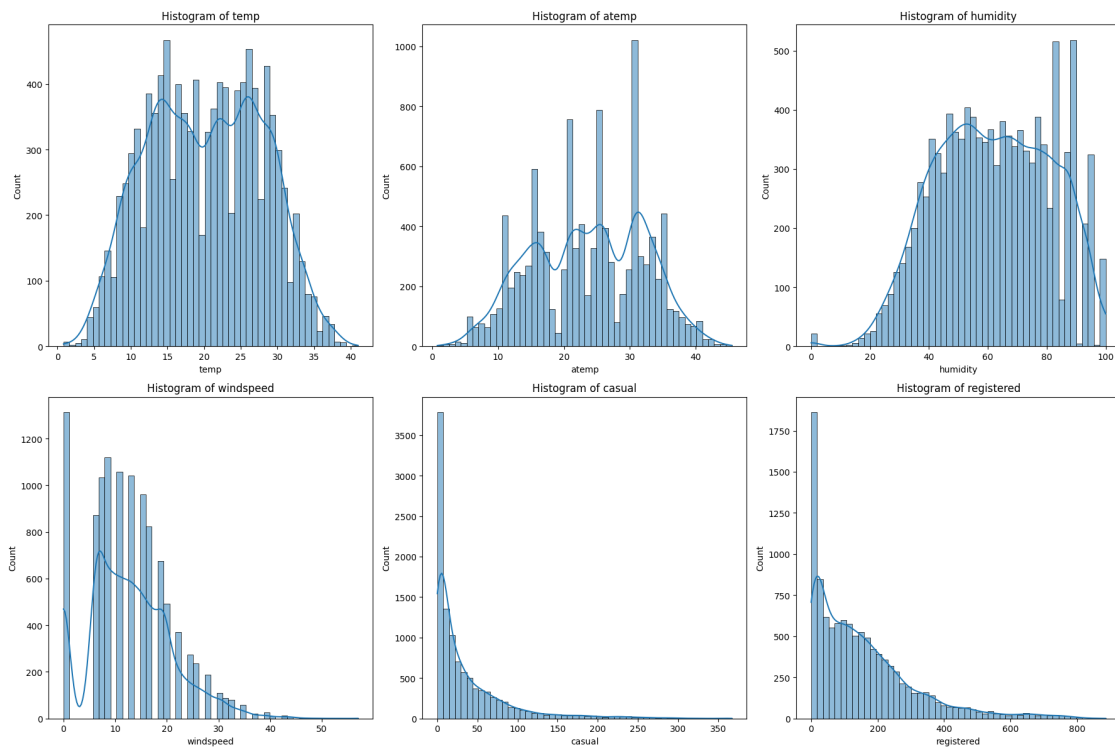
plt.subplot(2, 3, 4)
plot_histplot(df, 'windspeed', 234)

plt.subplot(2, 3, 5)
plot_histplot(df, 'casual', 235)

plt.subplot(2, 3, 6)
plot_histplot(df, 'registered', 236)

plt.tight_layout()
plt.show()

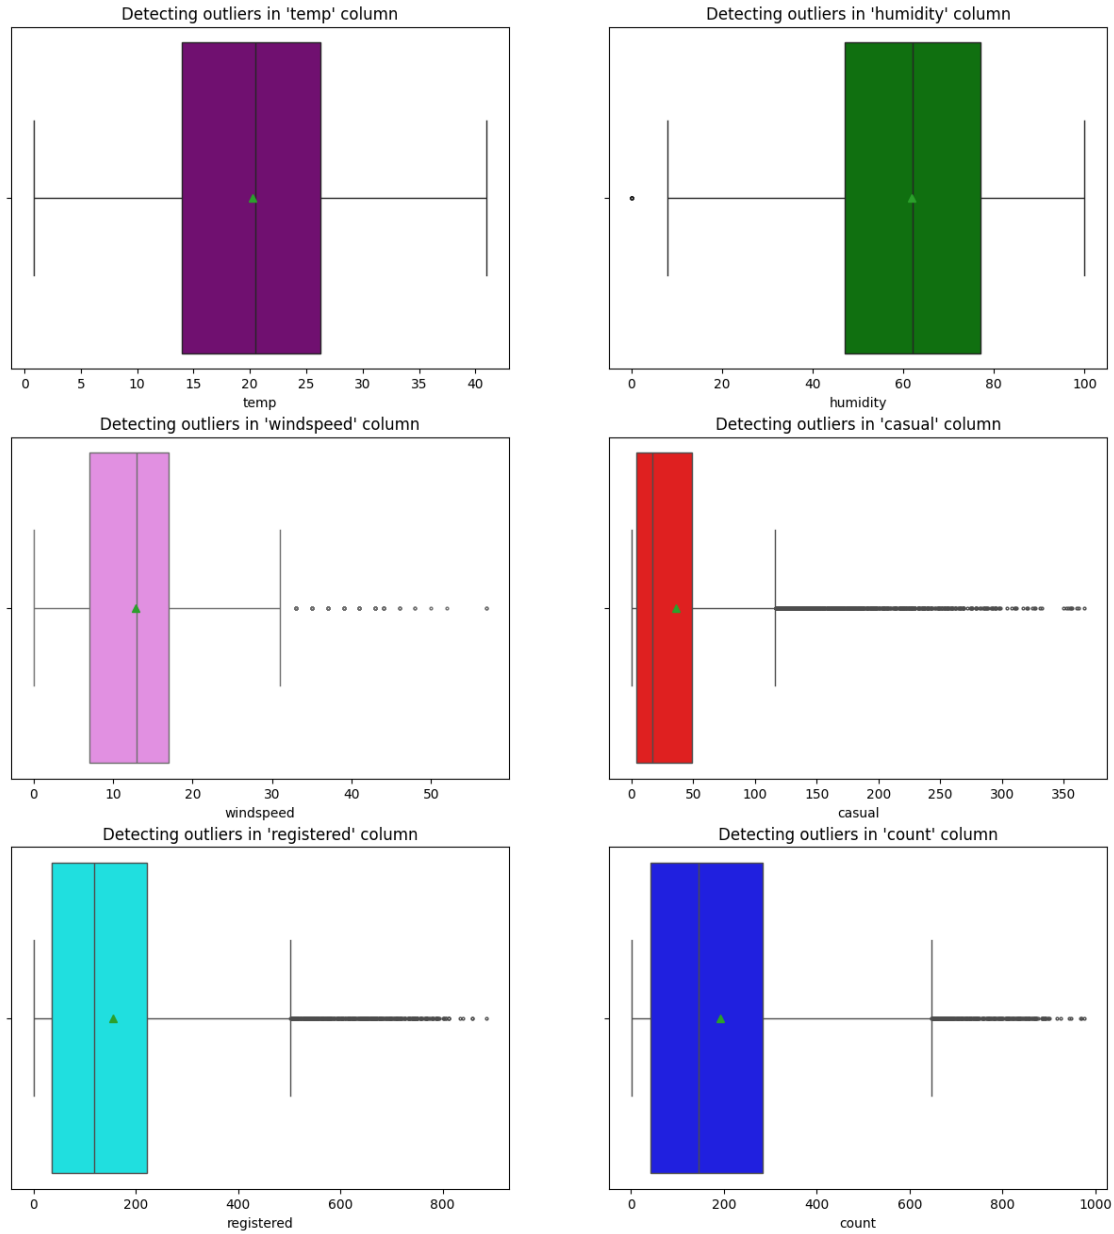
```



0.0.1 Detecting Outliers in the dataset

```
[58]: def plot_outliers(df, columns):
        colors = np.random.permutation(['green', 'blue', 'red', 'purple', 'cyan', 'violet'])
        count = 1
        plt.figure(figsize=(15, 16))
        for i in columns:
            plt.subplot(3, 2, count)
            plt.title(f"Detecting outliers in '{i}' column")
            sns.boxplot(data=df, x=df[i], color=colors[count - 1], showmeans=True, fliersize=2)
            plt.plot()
            count += 1

columns = ['temp', 'humidity', 'windspeed', 'casual', 'registered', 'count']
plot_outliers(df, columns)
```



1 Exploratory Analysis

- The lowest average hourly count of rental bikes is observed in January, followed by February and March.
- Out of every 100 users, approximately 19 are casual users, and 81 are registered users.
- Over 85% of the recorded windspeed data has a value of less than 20.
- The mean total hourly count of rental bikes is 144 for the year 2011 and 239 for the year 2012, indicating an annual growth rate of 65.41%.
- The dataset spans from January 1, 2011, to December 19, 2012, totaling 718 days and 23

hours.

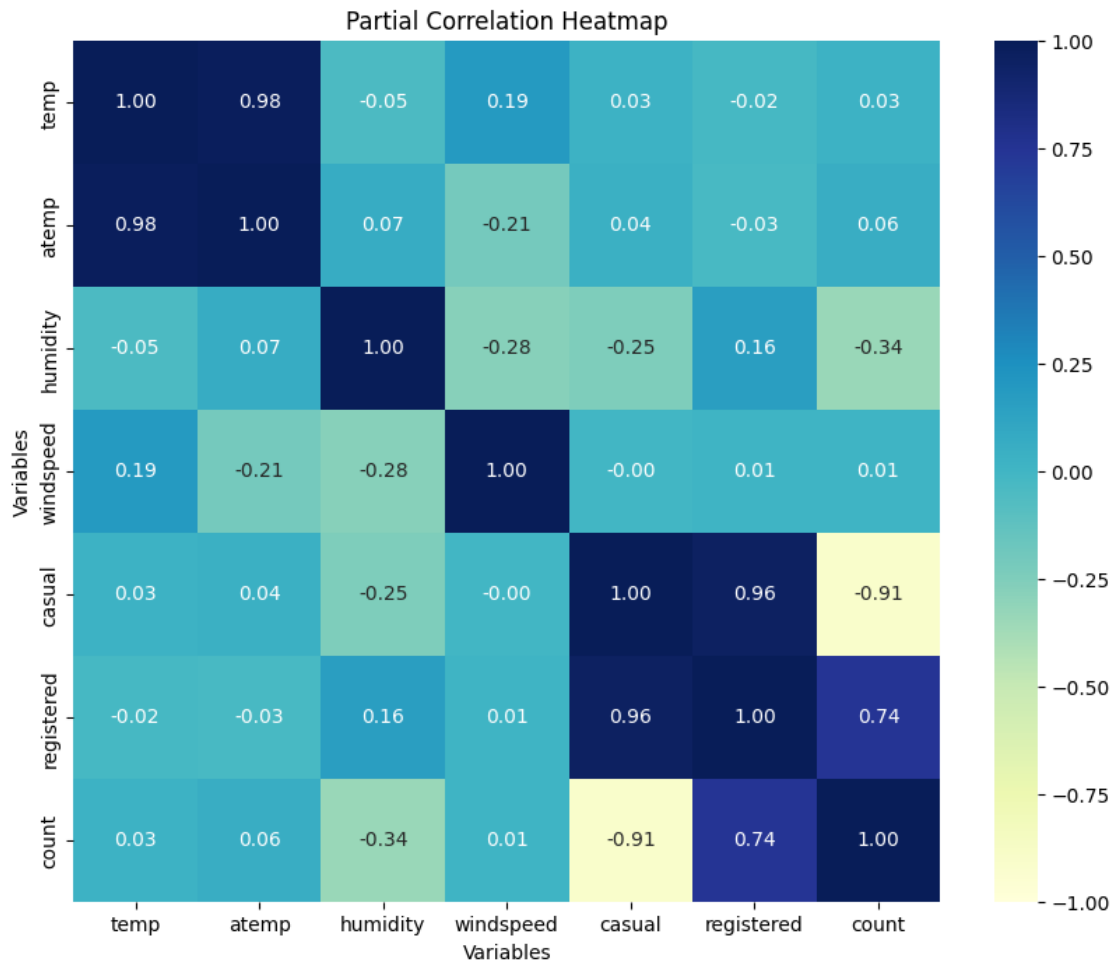
- More than 80% of the time, the temperature remains below 28 degrees Celsius.
- Similarly, more than 80% of the time, the humidity value exceeds 40, indicating varying levels from optimum to too moist.
- Throughout the day, there is a notable fluctuation in counts, with lower counts during early morning hours, a morning peak, a peak count in the afternoon, and a gradual decline in the evening and nighttime.
- The count of rental bikes exhibits a seasonal pattern, with higher demand during the season_1 and season_2 months, a slight decline in the season_3, and further decrease in season_4.

2 Correlation between the Variables

Partial Correlations between Variables

```
[59]: import pingouin as pg
def plot_partial_corr_heatmap(df):
    df_subset = df
    partial_corr = df_subset.pcorr()
    plt.figure(figsize=(10, 8))
    sns.heatmap(data=partial_corr, cmap='YlGnBu', annot=True, fmt=".2f",
    ↪vmin=-1, vmax=1)
    plt.title('Partial Correlation Heatmap')
    plt.xlabel('Variables')
    plt.ylabel('Variables')
    plt.show()


plot_partial_corr_heatmap(data)
```



Semi-Partial Correlations between Variables


```
[60]: def plot_semi_partial_corr_heatmap(df):
    new_df = df[['temp', 'atemp', 'humidity', 'windspeed', 'count']]
    corr_data = new_df.corr(method='spearman')
    plt.figure(figsize=(12, 8))
    sns.heatmap(data=corr_data, cmap='YlGnBu', annot=True, vmin=-1, vmax=1)
    plt.show()

plot_semi_partial_corr_heatmap(df)
```



```
[61]: def plot_correlation_heatmap(df):
    new_df = df[['temp', 'atemp', 'humidity', 'windspeed', 'count']]
    corr_data = new_df.corr()
    plt.figure(figsize=(12, 8))
    sns.heatmap(data=corr_data, cmap='YlGnBu', annot=True, vmin=-1, vmax=1)
    plt.show()

plot_correlation_heatmap(data)
```



2.0.1 Inferences on Correlations


- No high positive or negative correlations (0.7 - 0.9) are found between any columns.
- Low positive correlations (0.3 - 0.5) exist between the columns [count, temp], [count, atemp], and [casual, atemp].
- Moderate positive correlations (0.5 - 0.7) are observed between the columns [casual, count] and [casual, registered].
- A very high correlation (> 0.9) is observed between the columns [atemp, temp] and [count, registered].
- Negligible correlation is noted between all other combinations of columns

2.1 Does the presence of a working day influence the quantity of electric cycles rented ?

STEPS : Set up Null Hypothesis

```
[62]: def plot_boxplot(df, x_column, y_column):
    sns.boxplot(data=df, x=x_column, y=y_column)
    plt.title(f'Boxplot of {y_column} by {x_column}')
    plt.show()
```

```
plot_boxplot(df, 'workingday', 'count')
```



Visual examinations to ascertain whether the samples adhere to a normal distribution


```
[63]: def plot_workingday_comparison_hist(df, column, sample_size=2000):
    plt.figure(figsize=(15, 5))

    plt.subplot(1, 2, 1)
    sns.histplot(df.loc[df['workingday'] == 1, column].sample(sample_size),
                 element='step', color='purple', kde=True, label='workingday')
    plt.legend()

    plt.subplot(1, 2, 2)
    sns.histplot(df.loc[df['workingday'] == 0, column].sample(sample_size),
                 element='step', color='green', kde=True,
    label='non_workingday')
    plt.legend()

    plt.show()
```

```
plot_workingday_comparison_hist(df, 'count')
```



Based on the plot above, it can be deduced that the distributions deviate from the normal distribution.

Assessing distribution via QQ Plot


```
[64]: def qq_plot_workingday_comparison(df, column, sample_size=2000):
    plt.figure(figsize=(15, 6))
    plt.suptitle('QQ plots for the count of electric vehicles rented in_
    ↪workingday and non_workingday')

    plt.subplot(1, 2, 1)
    stats.probplot(df.loc[df['workingday'] == 1, column].sample(sample_size),
    ↪plot=plt, dist='norm')
    plt.title('QQ plot for workingday')

    plt.subplot(1, 2, 2)
    stats.probplot(df.loc[df['workingday'] == 0, column].sample(sample_size),
    ↪plot=plt, dist='norm')
    plt.title('QQ plot for non_workingday')

    plt.show()

qq_plot_workingday_comparison(df, 'count')
```



In a QQ plot, if data points closely align with the diagonal line, the distribution is likely normal.

Conducting the Shapiro-Wilk test to assess normality. H_0 : The sample conforms to a normal distribution.

H_1 : The sample deviates from a normal distribution.

Significance Level (α) = 0.05

Test Statistics: Shapiro-Wilk test for normality

```
[65]: def shapiro_test_normality(data, column, condition_column, condition_value,
    ↪sample_size=2000, alpha=0.05):
    sample = data.loc[data[condition_column] == condition_value, column].
    ↪sample(sample_size)
    test_stat, p_value = stats.shapiro(sample)
    print('p-value:', p_value)
    if p_value < alpha:
        print('The sample does not follow a normal distribution')
    else:
        print('The sample follows a normal distribution')

    print("Workingday")
    shapiro_test_normality(df, 'count', 'workingday', 1)
    print()
    print("Non-Workingday")
    shapiro_test_normality(df, 'count', 'workingday', 0)
```

Workingday

p-value: 2.914531643120015e-38

The sample does not follow a normal distribution

Non-Workingday
p-value: 9.117709753918228e-36
The sample does not follow a normal distribution

Applying the Box-Cox transformation to the data and assessing whether the transformed data adheres to a normal distribution.

```
[66]: def boxcox_shapiro_test_normality(data, column, condition_column,
    ↪condition_value, alpha=0.05):
    transformed_data = stats.boxcox(data.loc[data[condition_column] ==
    ↪condition_value, column])[0]
    test_stat, p_value = stats.shapiro(transformed_data)
    print('p-value:', p_value)
    if p_value < alpha:
        print('The sample does not follow a normal distribution')
    else:
        print('The sample follows a normal distribution')

    print("Workingday")
    boxcox_shapiro_test_normality(df, 'count', 'workingday', 1)
    print()
    print("Non-Workingday")
    boxcox_shapiro_test_normality(df, 'count', 'workingday', 0)
```

Workingday
p-value: 1.606449722752868e-33
The sample does not follow a normal distribution

Non-Workingday
p-value: 8.140929444965395e-24
The sample does not follow a normal distribution

/var/folders/kk/7w6727t942z6xwr_96jpcwtc0000gn/T/ipykernel_989/1579685415.py:3:
UserWarning: scipy.stats.shapiro: For N > 5000, computed p-value may not be
accurate. Current N is 7412.

```
test_stat, p_value = stats.shapiro(transformed_data)
```

Workingday: *The sample does not follow a normal distribution ($p < 0.05$)*

Non-Workingday: *Similarly, the sample does not follow a normal distribution ($p < 0.05$)*

Both samples fail the test for normality.

Despite applying the Box-Cox transformation to both the “workingday” and “non_workingday” datasets, neither conforms to a normal distribution.

As the samples do not exhibit a normal distribution, the application of the T-Test is inappropriate.

Ho : Mean number of electric cycles rented is the same for working and non-working days

Ha : Mean number of electric cycles rented is not the same for working and non-working days

Assuming a significance level of 0.05

Test statistics: Mann-Whitney U rank test for two independent samples

```
[67]: import scipy.stats as stats

def mann_whitney_test(df, column, condition_column, condition_value, alpha=0.05):
    test_stat, p_value = stats.mannwhitneyu(df.loc[df[condition_column] == condition_value, column],
                                             df.loc[df[condition_column] != condition_value, column])
    print('P-value:', p_value)
    if p_value < alpha:
        print('Mean number of electric cycles rented is not the same for working and non-working days')
    else:
        print('Mean number of electric cycles rented is the same for working and non-working days')

mann_whitney_test(df, 'count', 'workingday', 1)
```

P-value: 0.9679139953914079

Mean number of electric cycles rented is the same for working and non-working days


Hence, there is no statistically significant difference in the mean hourly count of total rental bikes between working and non-working days.

2.2 Does the presence of holidays affect the number of electric cycles rented?

STEPS : Set up Null Hypothesis

```
[68]: def plot_boxplot(df, x_column, y_column):
    sns.boxplot(data=df, x=x_column, y=y_column)
    plt.title(f'Boxplot of {y_column} by {x_column}')
    plt.show()

plot_boxplot(df, 'holiday', 'count')
```



Visual examinations to ascertain whether the samples adhere to a normal distribution

```
[69]: def plot_holiday_comparison_hist(df, column, sample_size=2000):
    holiday_sample = df.loc[df['holiday'] == 1, column]
    non_holiday_sample = df.loc[df['holiday'] == 0, column]

    if sample_size > len(holiday_sample):
        holiday_sample_size = len(holiday_sample)
    else:
        holiday_sample_size = sample_size

    if sample_size > len(non_holiday_sample):
        non_holiday_sample_size = len(non_holiday_sample)
    else:
        non_holiday_sample_size = sample_size


    plt.figure(figsize=(15, 5))
    plt.subplot(1, 2, 1)
    sns.histplot(holiday_sample.sample(holiday_sample_size, replace=False),
                  element='step', color='purple', kde=True, label='holiday')
```

```

plt.legend()
plt.subplot(1, 2, 2)
sns.histplot(non_holiday_sample.sample(non_holiday_sample_size,
↪replace=False),
              element='step', color='green', kde=True, label='non_holiday')
plt.legend()
plt.show()

plot_holiday_comparison_hist(df, 'count')

```



Based on the plot above, it can be deduced that the distributions deviate from the normal distribution.

Assessing distribution via QQ Plot

```

[70]: def plot_qq_holiday_comparison(df, column, sample_size=200):
    plt.figure(figsize=(15, 6))
    plt.suptitle('QQ plots for the count of electric vehicles rented in holiday_
↪and non_holiday')


    plt.subplot(1, 2, 1)
    stats.probplot(df.loc[df['holiday'] == 1, column].sample(sample_size),
↪plot=plt, dist='norm')
    plt.title('QQ plot for holiday')

    plt.subplot(1, 2, 2)
    stats.probplot(df.loc[df['holiday'] == 0, column].sample(sample_size),
↪plot=plt, dist='norm')
    plt.title('QQ plot for non_holiday')

    plt.show()

plot_qq_holiday_comparison(df, 'count')

```



In a QQ plot, if data points closely align with the diagonal line, the distribution is likely normal.

Conducting the Shapiro-Wilk test to assess normality. H_0 : The sample conforms to a normal distribution.

H_1 : The sample deviates from a normal distribution.

Significance Level (α) = 0.05

Test Statistics: Shapiro-Wilk test for normality

```
[71]: def shapiro_test_holiday(df, column, holiday_value, alpha=0.05, sample_size=200):
    sample = df.loc[df['holiday'] == holiday_value, column].sample(sample_size)
    test_stat, p_value = stats.shapiro(sample)
    print('p-value:', p_value)
    if p_value < alpha:
        print('The sample does not follow normal distribution')
    else:
        print('The sample follows normal distribution')

    print('Holiday')
    shapiro_test_holiday(df, 'count', 1)
    print()
    print('Non-Holiday')
    shapiro_test_holiday(df, 'count', 0)
```

Holiday

p-value: 7.312238915724563e-11

The sample does not follow normal distribution

Non-Holiday

p-value: 1.721475698571981e-12
The sample does not follow normal distribution

Applying the Box-Cox transformation to the data and assessing whether the transformed data adheres to a normal distribution.

```
[72]: def boxcox_shapiro_test_holiday(df, column, holiday_value, alpha=0.05):
        transformed_data = stats.boxcox(df.loc[df['holiday'] == holiday_value,
        column])[0]
        test_stat, p_value = stats.shapiro(transformed_data)
        print('p-value:', p_value)
        if p_value < alpha:
            print('The sample does not follow normal distribution')
        else:
            print('The sample follows normal distribution')

        print('Holiday')
        boxcox_shapiro_test_holiday(df, 'count', 1)
        print()
        print('Non-Holiday')
        boxcox_shapiro_test_holiday(df, 'count', 0)
```

Holiday
p-value: 2.134933458313291e-07
The sample does not follow normal distribution

Non-Holiday
p-value: 1.411562913878583e-36
The sample does not follow normal distribution

/var/folders/kk/7w6727t942z6xwr_96jpcwtc0000gn/T/ipykernel_989/764229625.py:3:
UserWarning: scipy.stats.shapiro: For N > 5000, computed p-value may not be
accurate. Current N is 10575.

```
test_stat, p_value = stats.shapiro(transformed_data)
```

Holiday: *The sample does not follow a normal distribution ($p < 0.05$)*

Non-Holiday: *Similarly, the sample does not follow a normal distribution ($p < 0.05$)*

Both samples fail the test for normality.

Despite employing the Box-Cox transformation on both the “holiday” and “non-holiday” datasets, the samples do not conform to a normal distribution.***

As the samples do not exhibit a normal distribution, the application of the T-Test is not appropriate.

Ho : Number of electric cycles rented is similar for holidays and non-holidays

Ha : Number of electric cycles rented is not similar for holidays and non-holidays days

Assuming significance level to be 0.05

Test statistics : Mann-Whitney U rank test for two independent samples

```
[73]: def mann_whitney_holiday_test(df, column, alpha=0.05, sample_size=200):
    test_stat, p_value = stats.mannwhitneyu(df.loc[df['holiday'] == 0, column].
    ↪sample(sample_size),
                                           df.loc[df['holiday'] == 1, column].
    ↪sample(sample_size))
    print('P-value:', p_value)
    if p_value < alpha:
        print('Number of electric cycles rented is not similar for holidays and
    ↪non-holidays days')
    else:
        print('Number of electric cycles rented is similar for holidays and
    ↪non-holidays')

mann_whitney_holiday_test(df, 'count')
```

P-value: 0.8355488880092926


Number of electric cycles rented is similar for holidays and non-holidays

Thus, the quantity of electric cycles rented shows statistical similarity between holidays and non-holidays.

2.3 Does the number of rented cycles remain consistent or vary across different weather conditions?

```
[74]: def plot_boxplot_weather(df):
    sns.boxplot(data=df, x='weather', y='count', showmeans=True)
    plt.show()

plot_boxplot_weather(df)
df_weather1 = df.loc[df['weather'] == 1]
print("len(df_weather1) = ", len(df_weather1))
df_weather2 = df.loc[df['weather'] == 2]
print("len(df_weather2) = ", len(df_weather2))
df_weather3 = df.loc[df['weather'] == 3]
print("len(df_weather3) = ", len(df_weather3))
df_weather4 = df.loc[df['weather'] == 4]
print("len(df_weather4) = ", len(df_weather4))
```



```
len(df_weather1) = 7192
len(df_weather2) = 2834
len(df_weather3) = 859
len(df_weather4) = 1
```

STEPS : Set up Null Hypothesis

Visual examinations to ascertain whether the samples adhere to a normal distribution

```
[75]: import seaborn as sns
import matplotlib.pyplot as plt


def plot_histplot_weather(df_weather1, df_weather2, df_weather3,
    sample_size=500):
    plt.figure(figsize=(15, 4))
    plt.subplot(1, 3, 1)
    sns.histplot(df_weather1.loc[:, 'count'].sample(sample_size), bins=40,
        element='step', color='purple', kde=True, label='weather1')
    plt.legend()
    plt.subplot(1, 3, 2)
    sns.histplot(df_weather2.loc[:, 'count'].sample(sample_size), bins=40,
        element='step', color='blue', kde=True, label='weather2')
```

```

plt.legend()
plt.subplot(1, 3, 3)
sns.histplot(df_weather3.loc[:, 'count'].sample(sample_size), bins=40,
             element='step', color='brown', kde=True, label='weather3')
plt.legend()
plt.show()

plot_histplot_weather(df_weather1, df_weather2, df_weather3)

```



Based on the plot above, it can be deduced that the distributions deviate from the normal distribution.

Assessing distribution via QQ Plot

```

[76]: def plot_qq_weather(df_weather1, df_weather2, df_weather3, sample_size=500):
    plt.figure(figsize=(18, 6))
    plt.suptitle('QQ plots for the count of electric vehicles rented in_
    ↪different weathers')

    plt.subplot(1, 3, 1)
    stats.probplot(df_weather1.loc[:, 'count'].sample(sample_size), plot=plt,
    ↪dist='norm')
    plt.title('QQ plot for weather_1')

    plt.subplot(1, 3, 2)
    stats.probplot(df_weather2.loc[:, 'count'].sample(sample_size), plot=plt,
    ↪dist='norm')
    plt.title('QQ plot for weather_2')


    plt.subplot(1, 3, 3)
    stats.probplot(df_weather3.loc[:, 'count'].sample(sample_size), plot=plt,
    ↪dist='norm')
    plt.title('QQ plot for weather_3')

    plt.show()

```



```
plot_qq_weather(df_weather1, df_weather2, df_weather3)
```



In a QQ plot, if data points closely align with the diagonal line, the distribution is likely normal.

Conducting the Shapiro-Wilk test to assess normality. H_0 : The sample conforms to a normal distribution.

H_1 : The sample deviates from a normal distribution.

Significance Level (α) = 0.05

Test Statistics: Shapiro-Wilk test for normality

```
[77]: def shapiro_test_weather(df_weather, sample_size=500):
    for i, df in enumerate(df_weather, start=1):
        test_stat, p_value = stats.shapiro(df.loc[:, 'count']).
        ↪sample(sample_size)
        print(f'Weather {i}:')
        print('P-value:', p_value)
        if p_value < 0.05:
            print('The sample does not follow normal distribution')
        else:
            print('The sample follows normal distribution')
        print()

shapiro_test_weather([df_weather1, df_weather2, df_weather3])
```

Weather 1:

P-value: 8.614116830276308e-19

The sample does not follow normal distribution

Weather 2:

P-value: 2.967472244980559e-20

The sample does not follow normal distribution

Weather 3:
P-value: 1.6592084087139088e-25
The sample does not follow normal distribution

Applying the Box-Cox transformation to the data and assessing whether the transformed data adheres to a normal distribution

```
[78]: def boxcox_shapiro_test_weather(df_weather, sample_size=500):
    for i, df in enumerate(df_weather, start=1):
        df_size = len(df.loc[:, 'count'])
        if sample_size > df_size:
            sample_size = df_size
            print(f"Sample size reduced to {sample_size} due to population size_
↳limitation.")

        transformed_data = stats.boxcox(df.loc[:, 'count'].sample(sample_size,
↳replace=True))[0]
        test_stat, p_value = stats.shapiro(transformed_data)
        print(f'Weather {i}:')
        print('P-value:', p_value)
        if p_value < 0.05:
            print('The sample does not follow normal distribution')
        else:
            print('The sample follows normal distribution')
        print()

    boxcox_shapiro_test_weather([df_weather1, df_weather2, df_weather3])
```

Weather 1:
P-value: 4.7547315711167956e-08
The sample does not follow normal distribution

Weather 2:
P-value: 8.312635723806701e-07
The sample does not follow normal distribution

Weather 3:
P-value: 6.832002469702163e-05
The sample does not follow normal distribution

Weather 1: *The sample does not follow a normal distribution ($p < 0.05$)*

Weather 2: *Similarly, the sample does not follow a normal distribution ($p < 0.05$)*

Weather 3: *Likewise, the sample does not follow a normal distribution ($p < 0.05$)*

All weather samples fail the test for normality.

Due to the samples' lack of normal distribution and unequal variance, the `f_oneway` test cannot

Ho : Mean no. of cycles rented is same for different weather

Ha : Mean no. of cycles rented is different for different weather

Assuming significance Level to be 0.05

```
[79]: import numpy as np
import scipy.stats as stats

def kruskal_wallis_test_weather(df_weather1, df_weather2, df_weather3, alpha=0.
    ↪05):
    test_stat, p_value = stats.kruskal(df_weather1, df_weather2, df_weather3)
    print('Test Statistic =', test_stat)
    print('p value =', p_value)
    p_value = np.mean(p_value)
    if p_value < alpha:
        print('Reject Null Hypothesis')
    else:
        print('Failed to reject Null Hypothesis')

kruskal_wallis_test_weather(df_weather1, df_weather2, df_weather3)
```

```
Test Statistic = [1.36471292e+01 1.83091584e+00 5.37649760e+00 1.56915686e+01
1.08840000e+04 3.70017441e+01 4.14298489e+01 1.83168690e+03
2.80380482e+01 2.84639685e+02 1.73745440e+02 2.04955668e+02]
p value = [1.08783632e-03 4.00333264e-01 6.79999165e-02 3.91398508e-04
0.00000000e+00 9.22939752e-09 1.00837627e-09 0.00000000e+00
8.15859150e-07 1.55338046e-62 1.86920588e-38 3.12206618e-45]
Reject Null Hypothesis
```

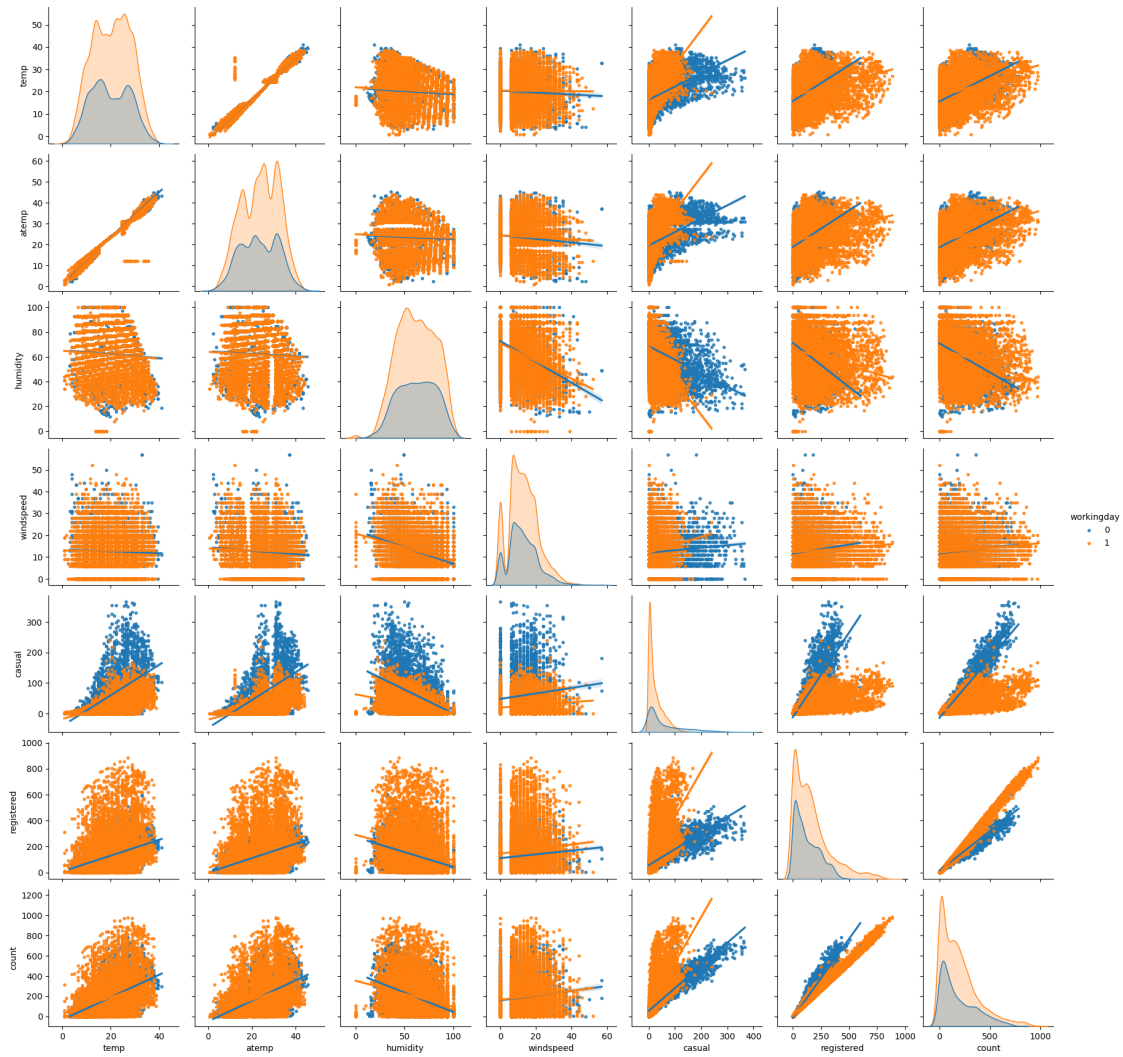
Hence, there is a statistically significant difference in the average number of rental bikes across varying weather conditions.

2.4 Does the number of rented cycles vary across different seasons ?

STEPS : Set up Null Hypothesis


```
[80]: def plot_pairplot_with_regression(df, hue_column='workingday'):
    sns.pairplot(data=df, kind='reg', hue=hue_column, markers='.')
    plt.show()

plot_pairplot_with_regression(df)
```



```
[81]: def boxplot_season_count(df):
    sns.boxplot(data=df, x='season', y='count', showmeans=True)
    plt.show()
```

```
boxplot_season_count(df)
df_season_1 = df.loc[df['season'] == 'season_1', 'count']
print("len(df_season_1) = ", len(df_season_1))
df_season_2 = df.loc[df['season'] == 'season_2', 'count']
print("len(df_season_2) = ", len(df_season_2))
df_season_3 = df.loc[df['season'] == 'season_3', 'count']
print("len(df_season_3) = ", len(df_season_3))
df_season_4 = df.loc[df['season'] == 'season_4', 'count']
print("len(df_season_4) = ", len(df_season_4))
```



```
len(df_season_1) = 2686
len(df_season_2) = 2733
len(df_season_3) = 2733
len(df_season_4) = 2734
```

Visual examinations to ascertain whether the samples adhere to a normal distribution

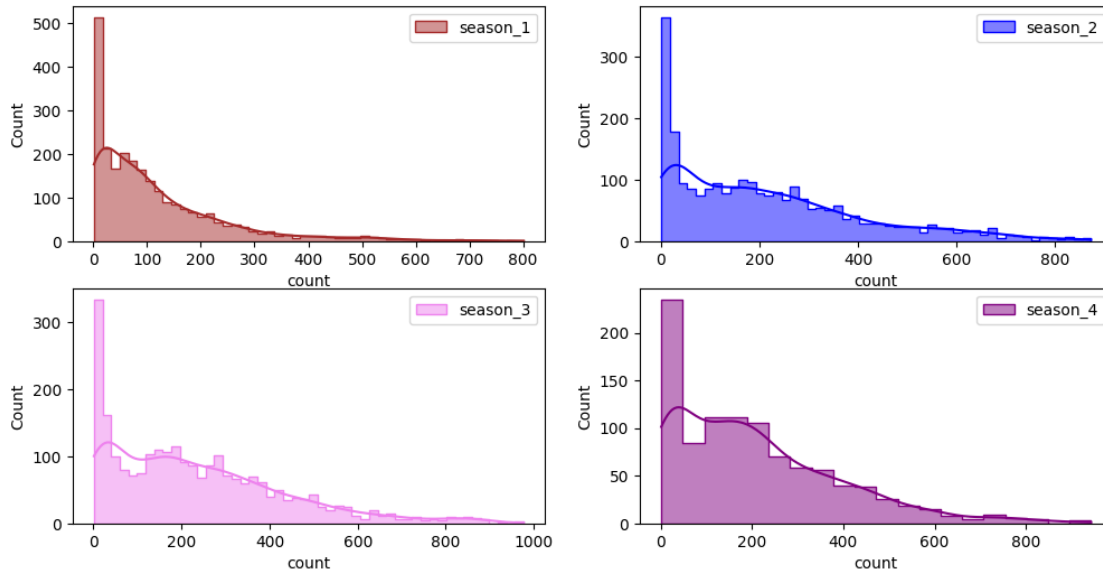
```
[82]: def plot_season_histograms(df_season_1, df_season_2, df_season_3, df_season_4):
    plt.figure(figsize=(12, 6))
    plt.subplot(2, 2, 1)
    sns.histplot(df_season_1.sample(2500), bins=50, element='step',
    color='brown', kde=True, label='season_1')
    plt.legend()
    plt.subplot(2, 2, 2)
    sns.histplot(df_season_2.sample(2500), bins=50, element='step',
    color='blue', kde=True, label='season_2')
    plt.legend()
    plt.subplot(2, 2, 3)
    sns.histplot(df_season_3.sample(2500), bins=50, element='step',
    color='violet', kde=True, label='season_3')
    plt.legend()
    plt.subplot(2, 2, 4)
```

```

sns.histplot(df_season_4.sample(1000), bins=20, element='step',
color='purple', kde=True, label='season_4')
plt.legend()
plt.show()

plot_season_histograms(df_season_1, df_season_2, df_season_3, df_season_4)

```



Based on the plot above, it can be deduced that the distributions deviate from the normal distribution.

Assessing distribution via QQ Plot

```

[83]: import matplotlib.pyplot as plt
import scipy.stats as stats

def plot_qq_plots_seasons(df_season_1, df_season_2, df_season_3, df_season_4):
    plt.figure(figsize=(12, 12))
    plt.suptitle('QQ plots for the count of electric vehicles rented in
different seasons')
    sample_size = min(2500, len(df_season_1))
    plt.subplot(2, 2, 1)
    stats.probplot(df_season_1.sample(sample_size), plot=plt, dist='norm')
    plt.title('QQ plot for season_1')
    sample_size = min(2500, len(df_season_2))
    plt.subplot(2, 2, 2)
    stats.probplot(df_season_2.sample(sample_size), plot=plt, dist='norm')
    plt.title('QQ plot for season_2')
    sample_size = min(2500, len(df_season_3))
    plt.subplot(2, 2, 3)

```

```

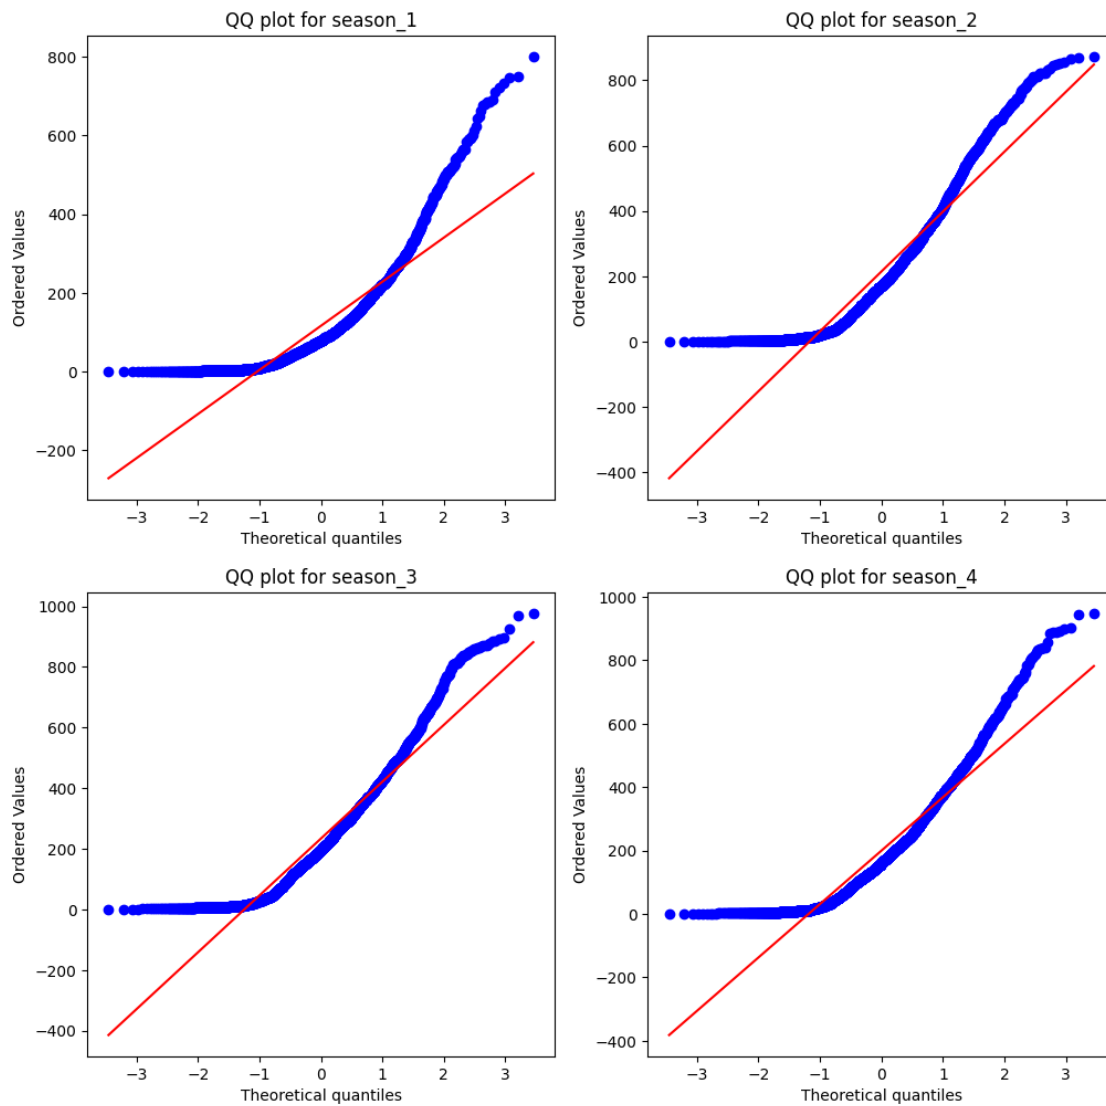
stats.probplot(df_season_3.sample(sample_size), plot=plt, dist='norm')
plt.title('QQ plot for season_3')
sample_size = min(2500, len(df_season_4))
plt.subplot(2, 2, 4)
stats.probplot(df_season_4.sample(sample_size), plot=plt, dist='norm')
plt.title('QQ plot for season_4')

plt.show()

plot_qq_plots_seasons(df_season_1, df_season_2, df_season_3, df_season_4)

```

QQ plots for the count of electric vehicles rented in different seasons



In a QQ plot, if data points closely align with the diagonal line, the distribution is likely normal.

Conducting the Shapiro-Wilk test to assess normality.

H_0 : The sample conforms to a normal distribution.

H_1 : The sample deviates from a normal distribution.

Significance Level (α) = 0.05

Test Statistics: Shapiro-Wilk test for normality

```
[84]: def shapiro_test_seasons(df_season_1, df_season_2, df_season_3, df_season_4):
    test_stat, p_value = stats.shapiro(df_season_1.sample(2500))
    print('Season_1:')
    print('p-value:', p_value)
    if p_value < 0.05:
        print('The sample does not follow normal distribution')
    else:
        print('The sample follows normal distribution')

    test_stat, p_value = stats.shapiro(df_season_2.sample(2500))
    print('\nSeason_2:')
    print('p-value:', p_value)
    if p_value < 0.05:
        print('The sample does not follow normal distribution')
    else:
        print('The sample follows normal distribution')

    test_stat, p_value = stats.shapiro(df_season_3.sample(2500))
    print('\nSeason_3:')
    print('p-value:', p_value)
    if p_value < 0.05:
        print('The sample does not follow normal distribution')
    else:
        print('The sample follows normal distribution')

    test_stat, p_value = stats.shapiro(df_season_4.sample(2500))
    print('\nSeason_4:')
    print('p-value:', p_value)
    if p_value < 0.05:
        print('The sample does not follow normal distribution')
    else:
        print('The sample follows normal distribution')

shapiro_test_seasons(df_season_1, df_season_2, df_season_3, df_season_4)
```

Season_1:

p-value: 2.3486954673453084e-47

The sample does not follow normal distribution

Season_2:

p-value: 3.650035838253269e-37

The sample does not follow normal distribution

Season_3:

p-value: 2.9123312767678484e-35

The sample does not follow normal distribution

Season_4:

p-value: 5.72429242739374e-38

The sample does not follow normal distribution

Applying the Box-Cox transformation to the data and assessing whether the transformed data adheres to a normal distribution

```
[85]: def boxcox_shapiro_test_seasons(df_season_1, df_season_2, df_season_3,
    ↪df_season_4):
    transformed_df_season_1 = stats.boxcox(df_season_1.sample(2500))[0]
    test_stat, p_value = stats.shapiro(transformed_df_season_1)
    print('Season_1:')
    print('p-value:', p_value)
    if p_value < 0.05:
        print('The sample does not follow normal distribution')
    else:
        print('The sample follows normal distribution')

    transformed_df_season_2 = stats.boxcox(df_season_2.sample(2500))[0]
    test_stat, p_value = stats.shapiro(transformed_df_season_2)
    print('\nSeason_2:')
    print('p-value:', p_value)
    if p_value < 0.05:
        print('The sample does not follow normal distribution')
    else:
        print('The sample follows normal distribution')

    transformed_df_season_3 = stats.boxcox(df_season_3.sample(2500))[0]
    test_stat, p_value = stats.shapiro(transformed_df_season_3)
    print('\nSeason_3:')
    print('p-value:', p_value)
    if p_value < 0.05:
        print('The sample does not follow normal distribution')
    else:
        print('The sample follows normal distribution')

    transformed_df_season_4 = stats.boxcox(df_season_4.sample(2500))[0]
    test_stat, p_value = stats.shapiro(transformed_df_season_4)
```

```

print('\nSeason_4:')
print('p-value:', p_value)
if p_value < 0.05:
    print('The sample does not follow normal distribution')
else:
    print('The sample follows normal distribution')

```

```

boxcox_shapiro_test_seasons(df_season_1, df_season_2, df_season_3, df_season_4)

```

Season_1:
p-value: 1.2712345021107283e-16
The sample does not follow normal distribution

Season_2:
p-value: 2.811170815583834e-21
The sample does not follow normal distribution

Season_3:
p-value: 7.10322875709124e-21
The sample does not follow normal distribution

Season_4:
p-value: 6.495794591680147e-20
The sample does not follow normal distribution

Season 1: *The sample does not follow a normal distribution ($p < 0.05$)*

Season 2: *Similarly, the sample does not follow a normal distribution ($p < 0.05$)*

Season 3: *Likewise, the sample does not follow a normal distribution ($p < 0.05$)*

Season 4: *Similarly, the sample does not follow a normal distribution ($p < 0.05$)*

All samples fail the test for normality.

Due to the lack of normal distribution and unequal variance among the samples, the f_oneway test

Ho : Mean no. of cycles rented is same for different seasons

Ha : Mean no. of cycles rented is different for different seasons

Assuming significance Level to be 0.05

```

[86]: def kruskal_test(df1, df2, df3, df4):
        alpha = 0.05
        test_stat, p_value = stats.kruskal(df1, df2, df3, df4)
        print('Test Statistic =', test_stat)
        print('p value =', p_value)

        if p_value < alpha:

```

```
print('Reject Null Hypothesis')
else:
    print('Failed to reject Null Hypothesis')

kruskal_test(df_season_1, df_season_2, df_season_3, df_season_4)
```

Test Statistic = 699.6668548181988

p value = 2.479008372608633e-151

Reject Null Hypothesis

Hence, there is a statistically significant difference in the average number of rental bikes across different seasons.

2.4.1 Inferences from Analysis

- The average hourly count of rental bikes shows no significant difference between working and non-working days.
- There is no statistically significant relationship between weather types 1, 2, and 3 and seasons concerning the average hourly total number of bikes rented.
- The hourly total number of rental bikes significantly varies across different seasons.
- The hourly total number of rental bikes varies significantly across different weather conditions.